


Bayesian Estimation of Daily Return and Volatility Using MCMC

*Project-1
MA4740: Bayesian Statistic
(Under the guidance of Prof. Arunabha Majumdar)*



October 19, 2025

Contents

1	Introduction	2
2	Methodology	2
2.1	Data Processing	2
2.2	Bayesian Model	2
2.3	MCMC Estimation via Rejection Sampling	2
3	Implementation	4
3.1	Data Processing (C++)	4
3.2	Visualization (Python)	7
3.3	MCMC Implementation (Python)	8
4	Results	20
4.1	Price Trend Analysis	20
4.2	Return Distribution	21
4.3	Volatility Analysis	22
4.4	Posterior Distributions	23
4.5	Posterior Predictive Distribution	24
4.6	Price Forecasts	25
5	Interpretation	26
5.1	Posterior Parameter Estimates	26
5.2	Predictive Distributions and Trading Implications	26
5.3	Risk Assessment and Position Sizing	26
5.4	Methodological Advantages	27
5.5	Economic and Financial Interpretation	27
5.6	Model Diagnostics and Reliability	27
6	Conclusion	28
7	Limitations	28

Abstract

This project applies Bayesian inference to XRP/USDT daily price data. Using a Normal-Inverse-Gamma conjugate prior model and Rejection Sampling, we estimate posterior distributions for the mean return (μ) and volatility (σ), compute credible intervals, and visualize the posterior predictive distribution to analyze uncertainty in cryptocurrency price behavior.

1 Introduction

Cryptocurrency markets are characterized by high volatility and uncertainty. Traditional frequentist approaches provide point estimates but fail to quantify uncertainty adequately. This project employs Bayesian methods to provide a complete probabilistic assessment of XRP/USDT returns and volatility.

2 Methodology

2.1 Data Processing

We process historical XRP/USDT price data to compute log returns:

$$r_t = \log\left(\frac{P_t}{P_{t-1}}\right) \quad (1)$$

where P_t is the price at time t and r_t is the log return.

2.2 Bayesian Model

We assume the log returns follow a Normal distribution:

$$r_t \sim \mathcal{N}(\mu, \sigma^2) \quad (2)$$

We use conjugate priors:

$$\mu \sim \mathcal{N}(\mu_0, \sigma_0^2) \quad (3)$$

$$\sigma^2 \sim \text{Inv-Gamma}(\alpha_0, \beta_0) \quad (4)$$

2.3 MCMC Estimation via Rejection Sampling

Rejection sampling is a Monte Carlo method that generates samples from the posterior distribution by accepting or rejecting candidate samples from a proposal distribution.

Algorithm:

1. Define proposal distributions for μ and σ^2 centered on data statistics:

$$q(\mu) = \mathcal{N}\left(\bar{r}, \sqrt{\sigma_0^2 + \frac{s^2}{n}}\right) \quad (5)$$

$$q(\sigma^2) = \text{Inv-Gamma}\left(\frac{n}{2}, \frac{1}{2} \sum_{i=1}^n (r_i - \bar{r})^2\right) \quad (6)$$

where \bar{r} is the sample mean of returns and s^2 is the sample variance.

2. For each iteration:

- Draw candidate (μ^*, σ^{2*}) from proposal $q(\mu, \sigma^2)$
- Calculate unnormalized posterior:

$$p^*(\mu^*, \sigma^{2*} | \mathbf{r}) = \prod_{i=1}^n \mathcal{N}(r_i | \mu^*, \sigma^{2*}) \cdot p(\mu^*) \cdot p(\sigma^{2*}) \quad (7)$$

- Calculate proposal density: $q(\mu^*, \sigma^{2*})$
- Calculate acceptance ratio:

$$\alpha = \frac{p^*(\mu^*, \sigma^{2*} | \mathbf{r})}{M \cdot q(\mu^*, \sigma^{2*})} \quad (8)$$

- Accept candidate with probability $\min(1, \alpha)$, where uniform random $u \sim U(0, 1)$:

$$\text{Accept if } u < \alpha \quad (9)$$

3. Repeat until desired number of samples ($n_{\text{samples}} = 2000$) is obtained

where $M = 10$ is an envelope constant chosen to satisfy $p^*(\theta | \mathbf{r}) \leq M \cdot q(\theta)$ for all θ .

Implementation Details:

- Maximum attempts: 1,000,000 iterations
- Target samples: 2,000 independent draws
- Progress monitoring: Every 100,000 attempts
- The proposal distribution is designed to center around the maximum likelihood estimates, improving acceptance rates

Advantages:

- Provides exact samples from the posterior (not approximations)
- Simple to implement and understand
- No need for iterative conditioning or burn-in period
- Independent samples (no autocorrelation between samples)
- Does not require tuning of step sizes or proposal variances

Disadvantages:

- Can have low acceptance rates (typically 1-20% for this problem)
- Requires many more proposals than the desired sample size
- Finding optimal envelope constant M requires careful consideration
- Computational cost increases with dimensionality
- May be inefficient in high-dimensional parameter spaces

3 Implementation

3.1 Data Processing (C++)

The data processing is implemented in C++ for efficiency:

Listing 1: Data Processing Code

```

1 #include<iostream>
2 #include<fstream>
3 #include<vector>
4 #include<string>
5 #include<sstream>
6 #include<algorithm>
7 #include<cmath>
8 #include<numeric>
9
10 using namespace std;
11
12 struct DailyData {
13     string date;
14     double price;
15 };
16
17 // Function to calculate skewness
18 double calculate_skewness(const vector<double>& data, double mean, double stddev) {
19     if (stddev == 0) return 0;
20     double skew = 0.0;
21     for (double val : data) {
22         skew += pow((val - mean) / stddev, 3);
23     }
24     return skew / data.size();
25 }
26
27 // Function to calculate kurtosis
28 double calculate_kurtosis(const vector<double>& data, double mean, double stddev) {
29     if (stddev == 0) return 0;
30     double kurt = 0.0;
31     for (double val : data) {
32         kurt += pow((val - mean) / stddev, 4);
33     }
34     return (kurt / data.size()) - 3.0; // Excess kurtosis
35 }
36
37 int main(){
38
39     ifstream data("data.csv");
40     if (!data.is_open()) {
41         cerr << "Error opening data.csv" << endl;
42         return 1;
43     }

```

```
44
45     vector<DailyData> records;
46     string line;
47
48     // Skip header lines
49     getline(data, line);
50     getline(data, line);
51
52     while(getline(data, line)) {
53         stringstream ss(line);
54         string field;
55
56         // leading comma
57         getline(ss, field, ',');
58         if (field.empty() && ss.eof()) continue;
59
60         DailyData record;
61
62         // Date
63         getline(ss, field, '"'); // consume until first quote
64         getline(ss, record.date, '"'); // read date
65         getline(ss, field, ','); // consume comma after date
66
67         // Price
68         getline(ss, field, ',');
69         try {
70             record.price = stod(field);
71             records.push_back(record);
72         } catch (const std::invalid_argument& ia) {
73             // Ignore lines that can't be parsed
74         }
75     }
76     data.close();
77
78     // Data is in reverse chronological order, so reverse it
79     reverse(records.begin(), records.end());
80
81     // --- Price Trend ---
82     ofstream prices_out("prices.csv");
83     prices_out << "Date,Price\n";
84     for(const auto& record : records) {
85         prices_out << '"' << record.date << "\",\"" << record.price << "\n";
86     }
87     prices_out.close();
88
89     // --- Returns Computation ---
90     vector<double> log_returns;
91     for(size_t i = 1; i < records.size(); ++i) {
92         if (records[i-1].price > 0) {
93             log_returns.push_back(log(records[i].price / records[i-1].price));
```

```
94     }
95 }
96
97 ofstream returns_out("returns.csv");
98 returns_out << "LogReturn\n";
99 for(double ret : log_returns) {
100     returns_out << ret << "\n";
101 }
102 returns_out.close();
103
104 // Calculate stats for log returns
105 double sum = accumulate(log_returns.begin(), log_returns.end(), 0.0);
106 double mean = sum / log_returns.size();
107 double sq_sum = inner_product(log_returns.begin(), log_returns.end(), log_returns.
108     begin(), 0.0);
109 double variance = (sq_sum / log_returns.size()) - mean * mean;
110 double stddev = sqrt(variance);
111 double skewness = calculate_skewness(log_returns, mean, stddev);
112 double kurtosis = calculate_kurtosis(log_returns, mean, stddev);
113
114 cout << "Log Returns Statistics:" << endl;
115 cout << "Mean: " << mean << endl;
116 cout << "Variance: " << variance << endl;
117 cout << "Skewness: " << skewness << endl;
118 cout << "Kurtosis: " << kurtosis << endl;
119
120 // --- Volatility Analysis (Rolling Mean & StdDev) ---
121 int window_size = 20;
122 ofstream rolling_out("rolling_stats.csv");
123 rolling_out << "Date,Price,RollingMean,RollingStd\n";
124 for(size_t i = 0; i < records.size(); ++i) {
125     rolling_out << ',' << records[i].date << "\n," << records[i].price;
126     if (i >= window_size - 1) {
127         double rolling_sum = 0.0;
128         for(int j = 0; j < window_size; ++j) {
129             rolling_sum += records[i-j].price;
130         }
131         double rolling_mean = rolling_sum / window_size;
132
133         double rolling_sq_sum = 0.0;
134         for(int j = 0; j < window_size; ++j) {
135             rolling_sq_sum += pow(records[i-j].price - rolling_mean, 2);
136         }
137         double rolling_std = sqrt(rolling_sq_sum / window_size);
138         rolling_out << "," << rolling_mean << "," << rolling_std;
139     } else {
140         rolling_out << ",,"; // No value for first entries
141     }
142     rolling_out << "\n";
143 }
```

```

143     rolling_out.close();
144
145     return 0;
146 }
```

3.2 Visualization (Python)

Visualization routines are implemented in Python:

Listing 2: Plotting Code

```

1 import pandas as pd
2 import matplotlib.pyplot as plt
3 import matplotlib.dates as mdates
4 import os
5
6 def ensure_images_dir():
7     """Creates the images directory if it doesn't exist."""
8     if not os.path.exists('images'):
9         os.makedirs('images')
10    print("Created 'images' directory for storing plots")
11
12 def plot_price_trend():
13     """Plots price trend over time."""
14     ensure_images_dir()
15     df = pd.read_csv('prices.csv', parse_dates=['Date'])
16     plt.figure(figsize=(12, 6))
17     plt.plot(df['Date'], df['Price'])
18     plt.title('XRP Price Trend')
19     plt.xlabel('Date')
20     plt.ylabel('Price (USD)')
21     plt.grid(True)
22     plt.gca().xaxis.set_major_formatter(mdates.DateFormatter('%b %Y'))
23     plt.gca().xaxis.set_major_locator(mdates.MonthLocator(interval=2))
24     plt.gcf().autofmt_xdate()
25     plt.savefig('images/price_trend.png')
26     plt.close()
27     print("Generated price_trend.png")
28
29 def plot_returns_histogram():
30     """Plots histogram of log returns."""
31     ensure_images_dir()
32     df = pd.read_csv('returns.csv')
33     plt.figure(figsize=(10, 6))
34     df['LogReturn'].hist(bins=50, density=True, alpha=0.7)
35     plt.title('Histogram of Daily Log Returns')
36     plt.xlabel('Log Return')
37     plt.ylabel('Density')
38     plt.grid(True)
39     plt.savefig('images/returns_histogram.png')
```

```

40 plt.close()
41 print("Generated returns_histogram.png")
42
43 def plot_rolling_stats():
44     """Plots rolling mean and standard deviation."""
45     ensure_images_dir()
46     df = pd.read_csv('rolling_stats.csv', parse_dates=['Date'])
47     plt.figure(figsize=(12, 8))
48
49     ax1 = plt.subplot(2, 1, 1)
50     plt.plot(df['Date'], df['Price'], label='Price', alpha=0.5)
51     plt.plot(df['Date'], df['RollingMean'], label='20-Day Rolling Mean', color='orange')
52     plt.title('Price and 20-Day Rolling Mean')
53     plt.ylabel('Price (USD)')
54     plt.legend()
55     plt.grid(True)
56     plt.setp(ax1.get_xticklabels(), visible=False)
57
58     plt.subplot(2, 1, 2, sharex=ax1)
59     plt.plot(df['Date'], df['RollingStd'], label='20-Day Rolling Std Dev', color='green')
60
61     plt.title('20-Day Rolling Standard Deviation (Volatility)')
62     plt.xlabel('Date')
63     plt.ylabel('Standard Deviation')
64     plt.legend()
65     plt.grid(True)
66
67     plt.gcf().autofmt_xdate()
68     plt.tight_layout()
69     plt.savefig('images/volatility_analysis.png')
70     plt.close()
71     print("Generated volatility_analysis.png")
72
73 if __name__ == '__main__':
74     plot_price_trend()
75     plot_returns_histogram()
76     plot_rolling_stats()

```

3.3 MCMC Implementation (Python)

The MCMC estimation and trading analysis:

Listing 3: MCMC and Analysis Code

```

1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 from scipy.stats import invgamma, norm
5 import os
6

```

```
7 def run_mcmc_estimation():
8     """
9     Performs Bayesian estimation of return mean and volatility using Rejection Sampling.
10    """
11    # Create images directory if it doesn't exist
12    if not os.path.exists('images'):
13        os.makedirs('images')
14        print("Created 'images' directory for storing plots")
15
16    # 1. Load the log returns data
17    try:
18        returns = pd.read_csv('returns.csv')['LogReturn'].dropna().to_numpy()
19    except FileNotFoundError:
20        print("Error: returns.csv not found. Please run the data_processor first.")
21        return
22
23    n = len(returns)
24    r_bar = np.mean(returns)
25    r_var = np.var(returns)
26
27    # 2. Set Priors (non-informative)
28    # For mu: Normal(mu_0, sigma_0^2)
29    mu_0 = 0.0
30    sigma_0_sq = 1000.0
31    # For sigma^2: Inv-Gamma(alpha_0, beta_0)
32    alpha_0 = 0.001
33    beta_0 = 0.001
34
35    # 3. Rejection Sampling setup
36    n_samples = 8000 # Target number of samples
37    max_attempts = 1000000 # Maximum attempts to avoid infinite loops
38
39    # Initialize storage for accepted samples
40    mu_samples = []
41    sigma_sq_samples = []
42
43    # 4. Run Rejection Sampling
44    print("Running Rejection Sampling...")
45    attempts = 0
46    accepted = 0
47
48    # Define proposal distributions (use simple Normal and Inv-Gamma)
49    # Proposal for mu: Normal centered at sample mean with larger variance
50    mu_proposal_mean = r_bar
51    mu_proposal_std = np.sqrt(r_var / n + sigma_0_sq)
52
53    # Proposal for sigma^2: Inverse-Gamma with parameters based on data
54    sigma_proposal_alpha = n / 2
55    sigma_proposal_beta = 0.5 * np.sum((returns - r_bar)**2)
```

```

57 # Calculate envelope constant M (upper bound on acceptance ratio)
58 # For simplicity, use a conservative value
59 M = 10.0
60
61 while accepted < n_samples and attempts < max_attempts:
62     attempts += 1
63
64     # Draw from proposal distributions
65     mu_prop = np.random.normal(mu_proposal_mean, mu_proposal_std)
66     sigma_sq_prop = invgamma.rvs(a=sigma_proposal_alpha, scale=sigma_proposal_beta)
67
68     # Calculate likelihood: product of Normal(mu, sigma^2) for each return
69     likelihood = np.prod(norm.pdf(returns, loc=mu_prop, scale=np.sqrt(sigma_sq_prop)))
70
71     # Calculate prior
72     prior_mu = norm.pdf(mu_prop, loc=mu_0, scale=np.sqrt(sigma_0_sq))
73     prior_sigma_sq = invgamma.pdf(sigma_sq_prop, a=alpha_0, scale=beta_0)
74     prior = prior_mu * prior_sigma_sq
75
76     # Calculate proposal density
77     proposal_mu = norm.pdf(mu_prop, loc=mu_proposal_mean, scale=mu_proposal_std)
78     proposal_sigma_sq = invgamma.pdf(sigma_sq_prop, a=sigma_proposal_alpha, scale=
79         sigma_proposal_beta)
80     proposal = proposal_mu * proposal_sigma_sq
81
82     # Calculate unnormalized posterior
83     posterior = likelihood * prior
84
85     # Calculate acceptance ratio
86     if proposal > 0:
87         acceptance_ratio = posterior / (M * proposal)
88     else:
89         acceptance_ratio = 0
90
91     # Accept or reject
92     if np.random.uniform(0, 1) < acceptance_ratio:
93         mu_samples.append(mu_prop)
94         sigma_sq_samples.append(sigma_sq_prop)
95         accepted += 1
96
97     # Print progress every 100000 attempts
98     if attempts % 100000 == 0:
99         acceptance_rate = accepted / attempts * 100
100        print(f"Attempts: {attempts}, Accepted: {accepted}, Acceptance Rate: {
101            acceptance_rate:.2f}%")
102
103 if accepted < n_samples:
104     print(f"Warning: Only {accepted} samples accepted out of {n_samples} target after
105         {attempts} attempts")

```

```
103     print(f"Acceptance rate: {accepted/attempts*100:.2f}%")
104 else:
105     print(f"Successfully generated {n_samples} samples with acceptance rate: {accepted/attempts*100:.2f}%")
106
107 # Convert to numpy arrays
108 mu_posterior = np.array(mu_samples)
109 sigma_posterior = np.sqrt(np.array(sigma_sq_samples)) # Convert variance to std dev
110
111 # 5. Visualize posterior distributions
112 plt.figure(figsize=(14, 6))
113
114 # Posterior for mu
115 plt.subplot(1, 2, 1)
116 plt.hist(mu_posterior, bins=50, density=True, alpha=0.7, label='Posterior of ')
117 plt.title('Posterior Distribution of Mean Return ()')
118 plt.xlabel('Mean Return')
119 plt.ylabel('Density')
120 plt.grid(True)
121
122 # Posterior for sigma
123 plt.subplot(1, 2, 2)
124 plt.hist(sigma_posterior, bins=50, density=True, alpha=0.7, label='Posterior of ',
125           color='green')
126 plt.title('Posterior Distribution of Volatility ()')
127 plt.xlabel('Volatility (Std. Dev. of Returns)')
128 plt.ylabel('Density')
129 plt.grid(True)
130
131 plt.tight_layout()
132 plt.savefig('images/mcmc_posteriors.png')
133 plt.close()
134 print("Generated mcmc_posteriors.png")
135
136 # 6. Generate and visualize posterior predictive distribution
137 # Sample future returns based on the posterior distributions
138 posterior_predictive = np.random.normal(
139     loc=np.random.choice(mu_posterior, size=5000),
140     scale=np.random.choice(sigma_posterior, size=5000)
141 )
142
143 plt.figure(figsize=(10, 6))
144 plt.hist(posterior_predictive, bins=50, density=True, alpha=0.7, color='purple')
145 plt.title('Posterior Predictive Distribution of Future Return')
146 plt.xlabel('Predicted Return')
147 plt.ylabel('Density')
148 plt.grid(True)
149
150 # Add vertical line at 0 for reference
151 plt.axvline(x=0, color='red', linestyle='--', alpha=0.7)
```

```

151 # Add a 95% credible interval on the plot
152 pred_ci = np.percentile(posterior_predictive, [2.5, 97.5])
153 plt.axvline(x=pred_ci[0], color='blue', linestyle='--', alpha=0.5)
154 plt.axvline(x=pred_ci[1], color='blue', linestyle='--', alpha=0.5)
155 plt.annotate(f'95% CI: [{pred_ci[0]:.4f}, {pred_ci[1]:.4f}]',
156                 xy=(0.05, 0.92), xycoords='axes fraction', fontsize=10)
157
158 plt.savefig('images/posterior_predictive.png')
159 plt.close()
160 print("Generated posterior_predictive.png")
161
162 # 7. Report credible intervals
163 mu_ci = np.percentile(mu_posterior, [2.5, 97.5])
164 sigma_ci = np.percentile(sigma_posterior, [2.5, 97.5])
165
166 print("\n--- Rejection Sampling Results ---")
167 print(f"95% Credible Interval for Mean Return (mu): [{mu_ci[0]:.6f}, {mu_ci[1]:.6f}]")
168
169 print(f"95% Credible Interval for Volatility (sigma): [{sigma_ci[0]:.6f}, {sigma_ci[1]:.6f}]")
170
171 print(f"95% Credible Interval for Predictive Return: [{pred_ci[0]:.6f}, {pred_ci[1]:.6f}]")
172
173 print("\nInterpretation:")
174 print(f"- Expected daily return is approximately {np.mean(mu_posterior):.6f}")
175 print(f"- The average volatility is {np.mean(sigma_posterior):.6f}")
176 print(f"- The wide credible intervals reflect high uncertainty in cryptocurrency
177     returns")
178 print(f"- There is a {:.1f}% probability of a positive return on any given day".
179       format(
180             100 * np.mean(posterior_predictive > 0)))
181
182 # 8. Generate dynamic analysis.md file
183 generate_analysis_report(returns, mu_posterior, sigma_posterior,
184                         posterior_predictive,
185                         mu_ci, sigma_ci, pred_ci)
186
187 def generate_analysis_report(returns, mu_posterior, sigma_posterior,
188                             posterior_predictive,
189                             mu_ci, sigma_ci, pred_ci):
190     """
191     Generates a comprehensive analysis report based on MCMC results
192     """
193
194     # Try to determine the asset name from the plots.py file
195     asset_name = "Financial Asset"
196
197     try:
198         with open('plots.py', 'r') as f:
199             for line in f:
200                 if "plt.title(' in line and "Price Trend')" in line:
201                     start = line.find("plt.title('") + len("plt.title('")
202
203

```

```
194         end = line.find(" Price Trend'"))
195         if end > start:
196             asset_name = line[start:end]
197             break
198     except:
199         pass
200
201     # Calculate key statistics
202     mean_return = np.mean(mu_posterior)
203     mean_volatility = np.mean(sigma_posterior)
204     prob_positive = np.mean(posterior_predictive > 0) * 100
205
206     # Determine risk level
207     if mean_volatility > 0.1:
208         risk_level = "HIGH"
209     elif mean_volatility > 0.05:
210         risk_level = "MEDIUM"
211     else:
212         risk_level = "LOW"
213
214     # Determine trading signal
215     if prob_positive > 60:
216         signal = "BUY"
217         confidence = "HIGH" if mu_ci[0] > 0 else "MEDIUM"
218     elif prob_positive < 40:
219         signal = "SELL"
220         confidence = "HIGH" if mu_ci[1] < 0 else "MEDIUM"
221     else:
222         signal = "NEUTRAL"
223         confidence = "MEDIUM" if (mu_ci[1] - mu_ci[0]) < 0.01 else "LOW"
224
225     # Calculate position sizing recommendation
226     base_position = 1.0 # Base position size (100%)
227     target_volatility = 0.05 # Target volatility level
228     vol_adjustment = min(1.0, target_volatility / mean_volatility)
229     directional_confidence = (prob_positive - 50) * 0.02 # -1.0 to 1.0
230     position_size = base_position * vol_adjustment * (0.5 + (directional_confidence *
231         0.5))
232     position_size = max(0.1, min(1.0, position_size)) # Clamp between 10% and 100%
233
234     # Format the position size as a percentage
235     position_size_pct = position_size * 100
236
237     # Get recent price data
238     recent_price = None
239     try:
240         prices_df = pd.read_csv('prices.csv')
241         if len(prices_df) > 0:
242             recent_price = prices_df['Price'].iloc[-1]
243     except:
```

```
243     pass
244
245 # Generate price predictions for different time horizons
246 price_predictions = {}
247 if recent_price is not None:
248     # Next day prediction - simulate using one-day returns
249     next_day_returns = np.random.normal(
250         loc=np.random.choice(mu_posterior, size=10000),
251         scale=np.random.choice(sigma_posterior, size=10000)
252     )
253     next_day_prices = recent_price * np.exp(next_day_returns)
254     next_day_ci = np.percentile(next_day_prices, [2.5, 97.5])
255
256 # Next week prediction - simulate 5 trading days
257 next_week_returns = np.zeros(10000)
258 for i in range(5): # 5 trading days in a week
259     daily_returns = np.random.normal(
260         loc=np.random.choice(mu_posterior, size=10000),
261         scale=np.random.choice(sigma_posterior, size=10000)
262     )
263     next_week_returns += daily_returns
264 next_week_prices = recent_price * np.exp(next_week_returns)
265 next_week_ci = np.percentile(next_week_prices, [2.5, 97.5])
266
267 # Next month prediction - simulate 22 trading days
268 next_month_returns = np.zeros(10000)
269 for i in range(22): # ~22 trading days in a month
270     daily_returns = np.random.normal(
271         loc=np.random.choice(mu_posterior, size=10000),
272         scale=np.random.choice(sigma_posterior, size=10000)
273     )
274     next_month_returns += daily_returns
275 next_month_prices = recent_price * np.exp(next_month_returns)
276 next_month_ci = np.percentile(next_month_prices, [2.5, 97.5])
277
278 price_predictions = {
279     'next_day': {
280         'low': next_day_ci[0],
281         'high': next_day_ci[1],
282         'expected': recent_price * np.exp(mean_return)
283     },
284     'next_week': {
285         'low': next_week_ci[0],
286         'high': next_week_ci[1],
287         'expected': recent_price * np.exp(5 * mean_return)
288     },
289     'next_month': {
290         'low': next_month_ci[0],
291         'high': next_month_ci[1],
292         'expected': recent_price * np.exp(22 * mean_return)
293     }
294 }
```

```

293         }
294     }
295
296     # Create the analysis report
297     with open('analysis.md', 'w') as f:
298         f.write(f"# {asset_name} Trading Analysis Report\n\n")
299         f.write("## Executive Summary\n")
300         f.write(f"Based on Bayesian MCMC analysis of historical price data, the following
301             insights and recommendations are generated:\n\n")
302
302         f.write("### Key Findings\n")
303         f.write(f"- **Expected Daily Return:** {mean_return:.6f} ({'+' if mean_return > 0
304             else ''}{mean_return*100:.4f}%) \n")
304         f.write(f"- **Daily Volatility:** {mean_volatility:.6f} ({mean_volatility*100:.4f
305             }%) \n")
305         f.write(f"- **Probability of Positive Return:** {prob_positive:.1f}%\n")
306         f.write(f"- **Risk Level:** {risk_level}\n")
307         f.write(f"- **95% Credible Interval for Mean Return:** [{mu_ci[0]:.6f}, {mu_ci
308             [1]:.6f}]\n")
308         f.write(f"- **95% Credible Interval for Volatility:** [{sigma_ci[0]:.6f}, {
309             sigma_ci[1]:.6f}]\n")
309         f.write(f"- **95% Predictive Interval for Next-Day Return:** [{pred_ci[0]:.6f}, {
310             pred_ci[1]:.6f}]\n\n")
310
311     # Add price predictions if available
312     if price_predictions and recent_price is not None:
313         f.write("### Price Forecasts (95% Credible Intervals)\n")
314         f.write(f"Current Price: {recent_price:.6f}\n\n")
315
316         f.write(" | Time Horizon | Expected Price | Lower Bound | Upper Bound | Range
317             |\n")
317         f.write(" |-----|-----|-----|-----|-----|\n")
318
318     next_day = price_predictions['next_day']
319     f.write(f" | Next Day | {next_day['expected']:.6f} | {next_day['low']:.6f} | {
320         next_day['high']:.6f} | {next_day['high'] - next_day['low']:.6f} |\n")
321
322     next_week = price_predictions['next_week']
323     f.write(f" | Next Week | {next_week['expected']:.6f} | {next_week['low']:.6f}
324         | {next_week['high']:.6f} | {next_week['high'] - next_week['low']:.6f} |\n")
325
325     next_month = price_predictions['next_month']
326     f.write(f" | Next Month | {next_month['expected']:.6f} | {next_month['low']:.6
327         f} | {next_month['high']:.6f} | {next_month['high'] - next_month['low']:.6
327         f} |\n\n")
328
328     # Additional price forecast visualizations

```

```

329     plt.figure(figsize=(12, 8))
330
331     # Plot current price and forecasts
332     horizons = ['Current', 'Next Day', 'Next Week', 'Next Month']
333     expected_prices = [recent_price, next_day['expected'], next_week['expected'],
334                         next_month['expected']]
335     lower_bounds = [recent_price, next_day['low'], next_week['low'], next_month['
336                         low']]
337     upper_bounds = [recent_price, next_day['high'], next_week['high'], next_month[
338                         'high']]
339
340     x = range(len(horizons))
341     plt.plot(x, expected_prices, 'o-', color='blue', linewidth=2, label='Expected
342                 Price')
343     plt.fill_between(x, lower_bounds, upper_bounds, color='blue', alpha=0.2,
344                       label='95% Credible Interval')
345
346     plt.xlabel('Time Horizon')
347     plt.ylabel('Price')
348     plt.title(f'{asset_name} Price Forecast')
349     plt.xticks(x, horizons)
350     plt.grid(True)
351     plt.legend()
352
353     plt.tight_layout()
354     plt.savefig('images/price_forecast.png')
355     plt.close()
356     f.write("![Price Forecast](images/price_forecast.png)\n\n")
357
358     f.write("## Trading Recommendation\n\n")
359     f.write(f"## Signal: {signal} ({confidence} Confidence)\n\n")
360
361     if signal == "BUY":
362         f.write(f"**Recommendation:** Enter a long position with {position_size_pct
363                         :.1f}% of available capital.\n\n")
364         if recent_price:
365             stop_loss = recent_price * (1 - 1.5 * mean_volatility)
366             take_profit = recent_price * (1 + 2 * mean_volatility)
367             f.write(f"- **Entry Price:** {recent_price:.6f}\n")
368             f.write(f"- **Stop Loss:** {stop_loss:.6f} (approximately {1.5 *
369                         mean_volatility * 100:.1f}% below entry)\n")
370             f.write(f"- **Take Profit:** {take_profit:.6f} (approximately {2 *
371                         mean_volatility * 100:.1f}% above entry)\n\n")
372
373     elif signal == "SELL":
374         f.write(f"**Recommendation:** Enter a short position with {position_size_pct
375                         :.1f}% of available capital.\n\n")
376         if recent_price:
377             stop_loss = recent_price * (1 + 1.5 * mean_volatility)
378             take_profit = recent_price * (1 - 2 * mean_volatility)
379             f.write(f"- **Entry Price:** {recent_price:.6f}\n")

```

```

370         f.write(f"- **Stop Loss:** {stop_loss:.6f} (approximately {1.5 *  
            mean_volatility * 100:.1f}% above entry)\n")  
371         f.write(f"- **Take Profit:** {take_profit:.6f} (approximately {2 *  
            mean_volatility * 100:.1f}% below entry)\n\n")  
372     else: # NEUTRAL  
373         f.write(f"**Recommendation:** Hold current positions or consider a neutral  
            strategy.\n\n")  
374         f.write(f"- Consider allocating {position_size_pct/2:.1f}% to long positions  
            and {position_size_pct/2:.1f}% to short positions.\n")  
375         f.write(f"- Alternatively, await stronger directional signals before entering  
            new positions.\n\n")  
376  
377     # Additional trading strategy based on forecast  
378     if price_predictions and recent_price is not None:  
379         next_day = price_predictions['next_day']  
380         expected_move_pct = (next_day['expected'] - recent_price) / recent_price *  
            100  
381  
382         f.write("### Short-Term Strategy Based on Price Forecast\n\n")  
383         if expected_move_pct > 1.0:  
384             f.write(f"The expected price movement for tomorrow is strongly positive ({  
                expected_move_pct:.2f}%). Consider a more aggressive long position,  
                potentially using call options or leveraged products if appropriate  
                for your risk tolerance.\n\n")  
385         elif expected_move_pct < -1.0:  
386             f.write(f"The expected price movement for tomorrow is strongly negative ({  
                expected_move_pct:.2f}%). Consider a more aggressive short position,  
                potentially using put options or leveraged products if appropriate for  
                your risk tolerance.\n\n")  
387         else:  
388             f.write(f"The expected price movement for tomorrow is relatively small ({  
                expected_move_pct:.2f}%). Consider focusing on range-bound trading  
                strategies or accumulating positions at favorable prices within the  
                predicted range.\n\n")  
389  
390         range_width_pct = (next_day['high'] - next_day['low']) / recent_price * 100  
391         f.write(f"The predicted price range for tomorrow spans {range_width_pct:.2f}%  
            of the current price, which suggests {'significant' if range_width_pct >  
            5 else 'moderate' if range_width_pct > 2 else 'limited'} intraday trading  
            opportunities.\n\n")  
392  
393     # Add the rest of the detailed analysis  
394     f.write("## Detailed Analysis\n\n")  
395     f.write("### Return Distribution\n")  
396     f.write(f"The analysis of historical returns shows an expected daily return of {  
            mean_return:.6f} with a volatility of {mean_volatility:.6f}. ")  
397  
398     if mu_ci[0] < 0 < mu_ci[1]:  
399         f.write("The 95% credible interval for the mean return contains zero,  
            indicating uncertainty about the true direction of returns.\n\n")

```

```

400     elif mu_ci[0] > 0:
401         f.write("The 95% credible interval for the mean return is entirely positive,
402                 suggesting a reliable upward trend.\n\n")
403     else:
404         f.write("The 95% credible interval for the mean return is entirely negative,
405                 suggesting a reliable downward trend.\n\n")
406
407     f.write("### Volatility Analysis\n")
408     f.write(f"The estimated volatility of {mean_volatility:.6f} indicates ")
409     if mean_volatility > 0.1:
410         f.write("high levels of price fluctuation, typical of cryptocurrency markets.
411                 Proper risk management is essential.\n\n")
412     elif mean_volatility > 0.05:
413         f.write("moderate levels of price fluctuation. Standard risk management
414                 practices are advised.\n\n")
415     else:
416         f.write("relatively stable price behavior. Tighter stop-losses can be
417                 considered.\n\n")
418
419     f.write("### Prediction for Next Trading Day\n")
420     f.write(f"Based on the posterior predictive distribution, there is a {
421         prob_positive:.1f}% probability of a positive return on the next trading day.
422         ")
423     f.write(f"The 95% predictive interval for the next-day return is [{pred_ci[0]:.6f
424         }, {pred_ci[1]:.6f}].\n\n")
425
426     if price_predictions and recent_price is not None:
427         f.write("### Extended Time Horizon Predictions\n\n")
428
429         # Next week analysis
430         next_week = price_predictions['next_week']
431         next_week_expected_change = (next_week['expected'] - recent_price) /
432             recent_price * 100
433         f.write(f"**One-Week Outlook:** The expected price after one week is {
434             next_week['expected']:.6f}, representing a {'+' if
435             next_week_expected_change >= 0 else ''}{next_week_expected_change:.2f}%
436             change from the current price. ")
437         f.write(f"The 95% credible interval for the one-week price is [{next_week['
438             low']:.6f}, {next_week['high']:.6f}].\n\n")
439
440         # Next month analysis
441         next_month = price_predictions['next_month']
442         next_month_expected_change = (next_month['expected'] - recent_price) /
443             recent_price * 100
444         f.write(f"**One-Month Outlook:** The expected price after one month is {
445             next_month['expected']:.6f}, representing a {'+' if
446             next_month_expected_change >= 0 else ''}{next_month_expected_change:.2f}%
447             change from the current price. ")
448         f.write(f"The 95% credible interval for the one-month price is [{next_month['
449             low']:.6f}, {next_month['high']:.6f}].\n\n")

```

```
432
433     f.write("These extended forecasts become increasingly uncertain with time
434         horizon. The predictions incorporate both parameter uncertainty and random
435         market movements, resulting in wider intervals for longer horizons.\n\n")
436
437     f.write("### Risk Assessment\n")
438     f.write(f"The current risk level is assessed as {risk_level}, based on the
439             estimated volatility and uncertainty in return predictions. ")
440
441     if risk_level == "HIGH":
442         f.write("This suggests using smaller position sizes and wider stop-losses.\n\
443             n")
444     elif risk_level == "MEDIUM":
445         f.write("This suggests standard position sizing and stop-loss practices.\n\n")
446     else:
447         f.write("This suggests the potential for larger position sizes, but still
448             with appropriate risk controls.\n\n")
449
450     f.write("## Methodology\n")
451     f.write("This analysis uses Bayesian inference with Rejection Sampling to
452         estimate the posterior distributions of mean return and volatility. ")
453     f.write("The rejection sampling algorithm draws candidates from proposal
454         distributions and accepts them based on the ratio of the posterior to the
455         proposal density. ")
456     f.write("This method provides exact samples from the posterior distribution but
457         can be less efficient than Gibbs sampling for high-dimensional problems. ")
458     f.write("The posterior predictive distribution incorporates both parameter
459         uncertainty and inherent market randomness.\n\n")
460
461     f.write("For multi-period forecasts (weekly and monthly), the analysis simulates
462         multiple daily returns using random draws from the posterior predictive
463         distribution ")
464     f.write("and compounds them to generate price paths. The reported intervals
465         represent the 95% credible range of these simulated paths.\n\n")
466
467     f.write("## Limitations and Disclaimers\n")
468     f.write("1. This analysis is based solely on historical price data and does not
469         incorporate fundamental factors, news events, or market sentiment.\n")
470     f.write("2. Past performance is not indicative of future results. Financial
471         markets are complex systems subject to numerous influences.\n")
472     f.write("3. This report is generated automatically and should be used as one
473         input among many for trading decisions.\n")
474     f.write("4. The model assumes returns follow a relatively stable distribution,
475         which may not hold during market regime changes.\n")
476     f.write("5. Longer-term forecasts are subject to increasing uncertainty and
477         should be interpreted with appropriate caution.\n\n")
478
479     f.write("## Report Generation\n")
480     f.write(f"This analysis was generated automatically on {pd.Timestamp.now()}.")
```

```

463     strftime('%Y-%m-%d %H:%M:%S')} based on available historical data.")
464     print("Generated analysis.md with trading recommendations including price forecasts
465           for next day, week, and month")
466 if __name__ == '__main__':
467     run_mcmc_estimation()

```

4 Results

4.1 Price Trend Analysis

The figure below shows the historical price trend of XRP/USDT over the entire observation period. We can observe significant volatility and several distinct price movements, including periods of rapid appreciation and decline. This visual representation helps identify overall market trends and potential regime changes in the cryptocurrency's behavior.




Figure 1: XRP/USDT Price Trend Over Time

4.2 Return Distribution

The histogram of daily log returns reveals the distribution of price changes. The distribution appears approximately symmetric around zero with heavy tails, indicating that extreme price movements (both positive and negative) occur more frequently than would be expected under a normal distribution. This is a characteristic feature of cryptocurrency markets and justifies our Bayesian approach to quantify uncertainty.




Figure 2: Distribution of Daily Log Returns

4.3 Volatility Analysis

This figure displays two key metrics: the 20-day rolling mean (top panel) which smooths out short-term fluctuations to reveal underlying trends, and the 20-day rolling standard deviation (bottom panel) which measures volatility over time. We observe that volatility is not constant but varies significantly across different periods, with some periods showing much higher uncertainty than others. This time-varying volatility is crucial for risk management and position sizing decisions.




Figure 3: Rolling Mean and Volatility Analysis

4.4 Posterior Distributions

These histograms show the posterior distributions for the two key parameters estimated using rejection sampling. The left panel displays the posterior distribution of the mean return (μ), representing our belief about the average daily return after observing the data. The right panel shows the posterior distribution of volatility (σ), quantifying uncertainty in the daily standard deviation of returns. The shapes of these distributions capture both the parameter estimates and the uncertainty around them, which is the hallmark of Bayesian inference.




Figure 4: Posterior Distributions of Mean Return (μ) and Volatility (σ)

4.5 Posterior Predictive Distribution

The posterior predictive distribution represents our probabilistic forecast for the next day's return, incorporating both parameter uncertainty and inherent market randomness. The distribution is centered around the expected return, with the spread reflecting the combined effects of volatility and parameter uncertainty. The dashed blue lines indicate the 95% credible interval, showing the range within which we expect the next return to fall with 95% probability. The red vertical line at zero serves as a reference point, helping visualize the probability of positive versus negative returns.




Figure 5: Posterior Predictive Distribution for Future Returns

4.6 Price Forecasts

This visualization presents price forecasts for multiple time horizons: next day, next week (5 trading days), and next month (22 trading days). The blue line connects the expected prices, while the shaded region represents the 95% credible interval. Notice how the uncertainty (width of the credible interval) increases with the forecast horizon, reflecting the compounding effects of daily uncertainty over longer periods. This progressive widening of the credible interval is a natural consequence of the stochastic nature of returns and illustrates why longer-term predictions are inherently more uncertain.




Figure 6: Price Forecasts for Different Time Horizons

5 Interpretation

5.1 Posterior Parameter Estimates

The Bayesian analysis yields posterior distributions for both the mean return (μ) and volatility (σ) parameters. Based on the MCMC samples:

- **Mean Return (μ):** The posterior mean is approximately -0.0012 (or -0.12% daily), with a 95% credible interval of $[-0.088, 0.090]$. The credible interval contains zero, indicating substantial uncertainty about the true direction of expected returns. This suggests that the data does not provide strong evidence for a systematic upward or downward trend in XRP prices.
- **Volatility (σ):** The posterior mean is approximately 0.0545 (or 5.45% daily), with a 95% credible interval of $[0.053, 0.056]$. This relatively narrow interval indicates that volatility is estimated with higher precision than the mean return. The moderate-to-high volatility level is characteristic of cryptocurrency markets.

5.2 Predictive Distributions and Trading Implications

The posterior predictive distribution for next-day returns reveals a 49% probability of positive returns, essentially a coin flip. The 95% predictive interval spans from -14.3% to $+15.0\%$, illustrating the substantial day-to-day price uncertainty inherent in this asset class.

For extended time horizons:

- **One-Week Forecast:** Expected price decline of approximately 0.6% , with 95% credible interval spanning a potential 29% loss to 39% gain
- **One-Month Forecast:** Expected price decline of approximately 2.6% , with 95% credible interval spanning a potential 52% loss to 97% gain

The widening credible intervals over longer horizons reflect the compounding effect of daily uncertainty. This progressive expansion quantifies the increasing difficulty of making precise forecasts further into the future.

5.3 Risk Assessment and Position Sizing

The analysis classifies XRP as a **MEDIUM** risk asset based on its volatility profile. The estimated daily volatility of 5.45% implies:

- Approximately 68% of daily returns fall within $\pm 5.45\%$
- Approximately 95% of daily returns fall within $\pm 10.9\%$
- Extreme movements beyond $\pm 15\%$ occur but are relatively rare

Given the near-neutral probability of positive returns (49%) and wide credible interval for mean return, the analysis recommends a **NEUTRAL** trading stance with low confidence. The suggested position sizing of approximately 22.5% allocated equally to long and short positions (or a market-neutral strategy) reflects this uncertainty.

5.4 Methodological Advantages

The Bayesian approach provides several key advantages over frequentist methods:

- **Complete Uncertainty Quantification:** Rather than point estimates, we obtain full probability distributions that capture all aspects of parameter uncertainty
- **Coherent Probabilistic Framework:** All inferences are expressed as probabilities, facilitating decision-making under uncertainty
- **Exact Posterior Sampling:** Rejection sampling produces exact draws from the posterior distribution, without the approximation errors or convergence concerns of other MCMC methods
- **Statistical Independence:** Each posterior sample is independent, eliminating autocorrelation that can plague other sampling schemes
- **Integrated Prediction:** The posterior predictive distribution naturally combines parameter uncertainty with sampling variability, providing realistic prediction intervals
- **Interpretable Credible Intervals:** Unlike frequentist confidence intervals, Bayesian credible intervals have the intuitive interpretation that there is a 95% probability the true parameter lies within the interval given the data

5.5 Economic and Financial Interpretation

The negative expected return (-0.12% daily, approximately -30% annually if compounded) combined with high volatility suggests that XRP exhibits risk without commensurate expected reward over the sample period. However, the wide credible interval containing positive values indicates this negative drift is not statistically reliable.

The high volatility relative to the uncertain mean return implies:

- High Sharpe ratio uncertainty (risk-adjusted return is poorly determined)
- Substantial portfolio risk from position concentration
- Potential for both large gains and losses
- Need for active risk management and position monitoring

The substantial intraday price range (approximately 30% of current price) suggests opportunities for short-term trading strategies, though such strategies require careful consideration of transaction costs and execution risk.

5.6 Model Diagnostics and Reliability

The analysis successfully generated 2,000 independent posterior samples through rejection sampling, with typical acceptance rates of 1-20%. The posterior distributions exhibit:

- Smooth, unimodal shapes indicating convergence to the true posterior
- Reasonable dispersion reflecting appropriate uncertainty quantification

- Consistency with Maximum Likelihood estimates (which center the proposal distributions)

The wide credible intervals, rather than indicating model failure, accurately reflect the genuine uncertainty present in cryptocurrency price data. This honest assessment of uncertainty is a strength of the Bayesian approach.

6 Conclusion

This project demonstrates the application of Bayesian rejection sampling methods to cryptocurrency price analysis. The Normal-Inverse-Gamma model with rejection sampling provides robust estimates of return characteristics and uncertainty. While the acceptance rate is relatively low (typically 20-26%), rejection sampling offers the advantage of exact posterior samples without burn-in periods or convergence diagnostics. The generated trading analysis report offers actionable insights while acknowledging the limitations of purely statistical approaches.

The implementation successfully generates 2,000 posterior samples, which are then used to construct credible intervals and posterior predictive distributions. These results provide a complete probabilistic characterization of XRP/USDT price behavior.

7 Limitations

- Assumes returns are independent and identically distributed
- Does not incorporate market microstructure or fundamental factors
- Model may not capture regime changes or structural breaks
- Past performance does not guarantee future results
- Rejection sampling can be computationally intensive due to low acceptance rates (20-25%)
- The choice of envelope constant M affects efficiency but not correctness
- Non-informative priors may not fully represent prior knowledge